

Open Information Systems

Lecture 7: UPLIFT – Mapping Relational Databases to RDF

Christophe Debruyne

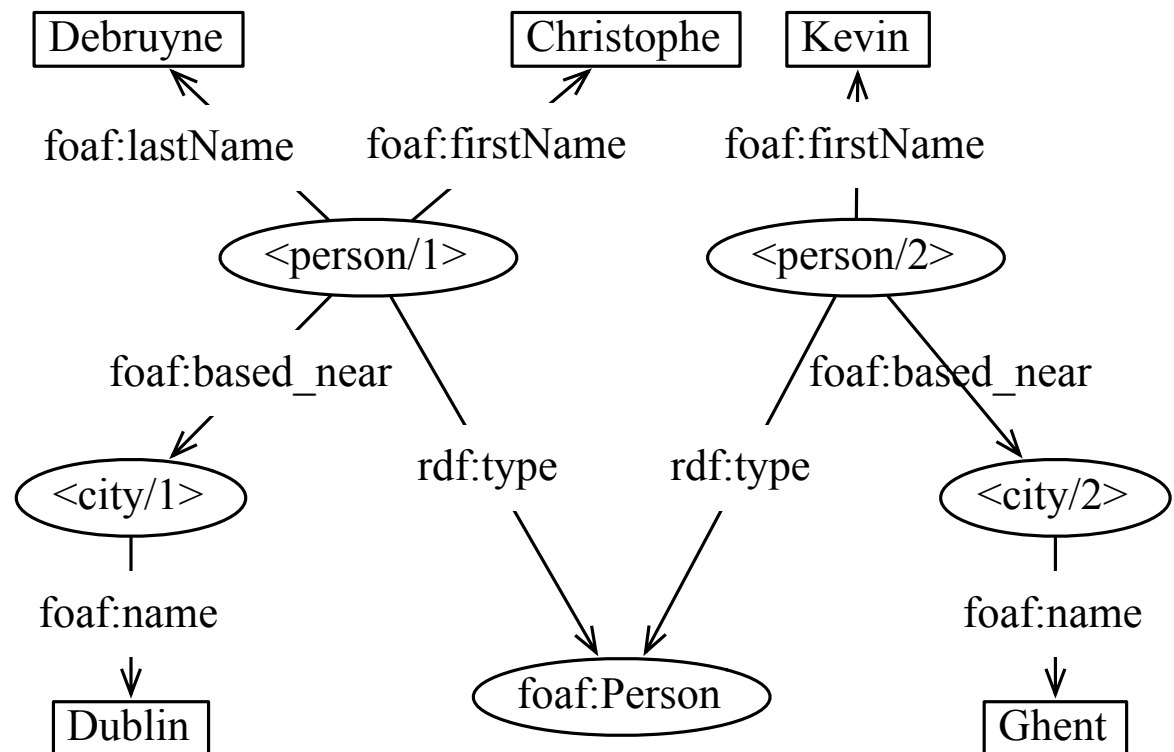
RDB2RDF

Person

ID	Fist	Last	CityID
1	Christophe	Debruyne	1
2	Kevin	NULL	2

City

ID	Name
1	Dublin
2	Ghent



RDB2RDF: W3C Recommendations

- It all started with Tim Berners-Lee proposing a direct mapping from RDBs to RDF.
- Over the years, two recommendations (standards) to map relational data to RDF emerged.
 - A **Direct Mapping** of relational data to RDF
 - **R2RML**: an RDB to RDF Mapping Language that is highly customizable to annotate relational data with ontologies to generate RDF.

RDB2RDF

The “started” started with Berners-Lee discussing a set of mappings – that can be generated automatically – between relational databases and RDF, see “Relational Databases on the Semantic Web” via <http://www.w3.org/DesignIssues/RDB-RDF.html>

“The semantic web data model is very directly connected with the model of relational databases. A relational database consists of tables, which consists of rows, or records. Each record consists of a set of fields. The record is nothing but the content of its fields, just as an RDF node is nothing but the connections: the property values. **The mapping is very direct**

- **a record is an RDF node;**
- **the field (column) name is RDF propertyType; and**
- **the record field (table cell) is a value.”**

Indeed, it will not always be that simple. Can you give examples?

Direct Mappings

TBL proposed a direct mapping. **Direct mappings** immediately reflect the structure of the database → The target RDF vocabulary directly reflects the names of database schema elements, and neither structure nor target vocabulary can be changed. [R2RML]

Process:

- **Existing table and column names are encoded into URIs.**
- Data is (i) *extracted*, (ii) *transformed* into RDF and then (iii) *loaded* into a triplestore.
- This is thus an ETL process.

This proposal – over time – was refined into a W3C recommendation, published in fall 2012, called “A Direct Mapping of Relational Data to RDF” <http://www.w3.org/TR/rdb-direct-mapping/>

Direct Mappings

- The database (both schema and data), primary keys and foreign keys are given to a *direct mapping engine* to produce an RDF graph.
 - Fields are mapping to literals;
 - Primary keys are used to construct URIs for resources;
 - And foreign keys are used to construct properties and relate resources.
- Example...

Direct Mapping Example

People		
PK		→Addresses(ID)
ID	fname	addr
1	Christophe	1
2	Kevin	NULL

Addresses	
PK	
ID	city
1	Brussels

@base <<http://foo.example/DB/>> .

@prefix xsd: <<http://www.w3.org/2001/XMLSchema#>> .

<People/ID=1> a <People> .

<People/ID=1> <People#ID> 1 .

<People/ID=1> <People#fname> "Christophe" .

<People/ID=1> <People#addr> 1 .

<People/ID=1> <People#ref-addr> <Addresses/ID=1> .

<People/ID=2> a <People> .

<People/ID=2> <People#ID> 2 .

<People/ID=2> <People#fname> "Kevin" .

<Addresses/ID=1> a <Addresses> .

<Addresses/ID=1> <Addresses#ID> 1 .

<Addresses/ID=1> <Addresses#city> "Brussels" .

Given a base URI

<http://foo.example/DB/>

Direct Mappings

Small discussion:

Are direct mappings meaningful?

Can you identify potential problems?

R2RML

- R2RML: RDB to RDF Mapping Language
 - A W3C Recommendation since fall 2012
 - <http://www.w3.org/TR/r2rml/>
- Creating an R2RML file that annotates a relational database with existing vocabularies and/or ontologies (RDFS or OWL).
- That R2RML file goes through an *R2RML Mapping Engine* to produce RDF.
- R2RML specified
 - An ontology to specify those mappings;
 - How those mappings should be interpreted to produce RDF.
 - R2RML files are thus stored as RDF.

R2RML: Running Example

Addresses	
PK	
ID	city
1	Brussels

What is being mapped? A logical table/view or an SQL query.

How to generate and state something about the subject of those triples.

How to generate predicates and objects.

```
@prefix rr: <http://www.w3.org/ns/r2rml#> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .
@prefix dbpedia: <http://dbpedia.org/ontology/> .
```

```
<#AddressTripleMap>
  a rr:TriplesMap ;
```

```
  rr:logicalTable [ rr:tableName "Addresses" ] ;
```

```
  rr:subjectMap [
    rr:template "http://foo.example/Addresses/{ID}" ;
    rr:class dbpedia:Place
  ] ;
```

```
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "city" ]
  ] ;
```

```
  .
```

R2RML: Running Example

Addresses	
PK	
ID	city
1	Brussels

```
@prefix rr: <http://www.w3.org/ns/r2rml#> .
@prefix foaf: <http://xmlns.com/foaf/0.1/> .
@prefix dbpedia: <http://dbpedia.org/ontology/> .

<#AddressTripleMap>
  a rr:TriplesMap ;
  rr:logicalTable [
    rr:sqlQuery """SELECT ID, city FROM
                  Addresses WHERE 1"""
  ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Addresses/{ID}" ;
    rr:class dbpedia:Place
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "city" ]
  ] ;
  .
```

R2RML: Running Example

People		
PK		→Addresses(ID)
ID	fname	addr
1	Christophe	1
2	Kevin	NULL

Addresses	
PK	
ID	city
1	Brussels

Relating People to
Addresses.

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ]
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [
        rr:child "addr" ;
        rr:parent "ID"
      ]
    ]
  ] ;
```

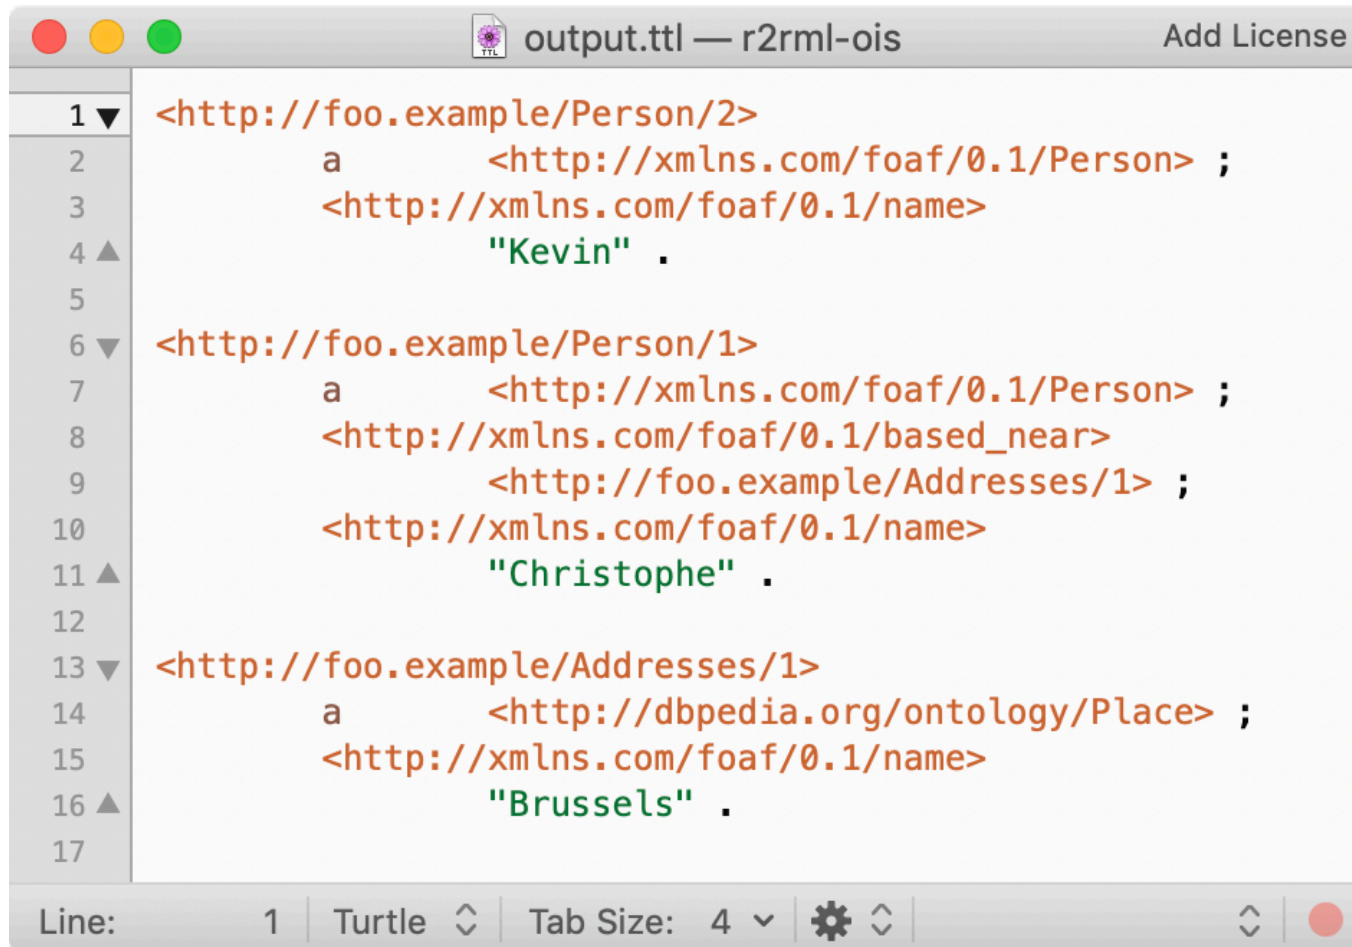
R2RML: Running Example

- Using R2RML-F (<http://github.com/chrdebru/r2rml>)
- The configuration file contains:

```
connectionURL = jdbc:mysql://localhost/r2rml
user = foo
password = bar
mappingFile = db-mapping.ttl
outputFile = output.ttl
format = TURTLE
```

- And then we execute:
\$ java -jar r2rml/r2rml.jar db-config.properties

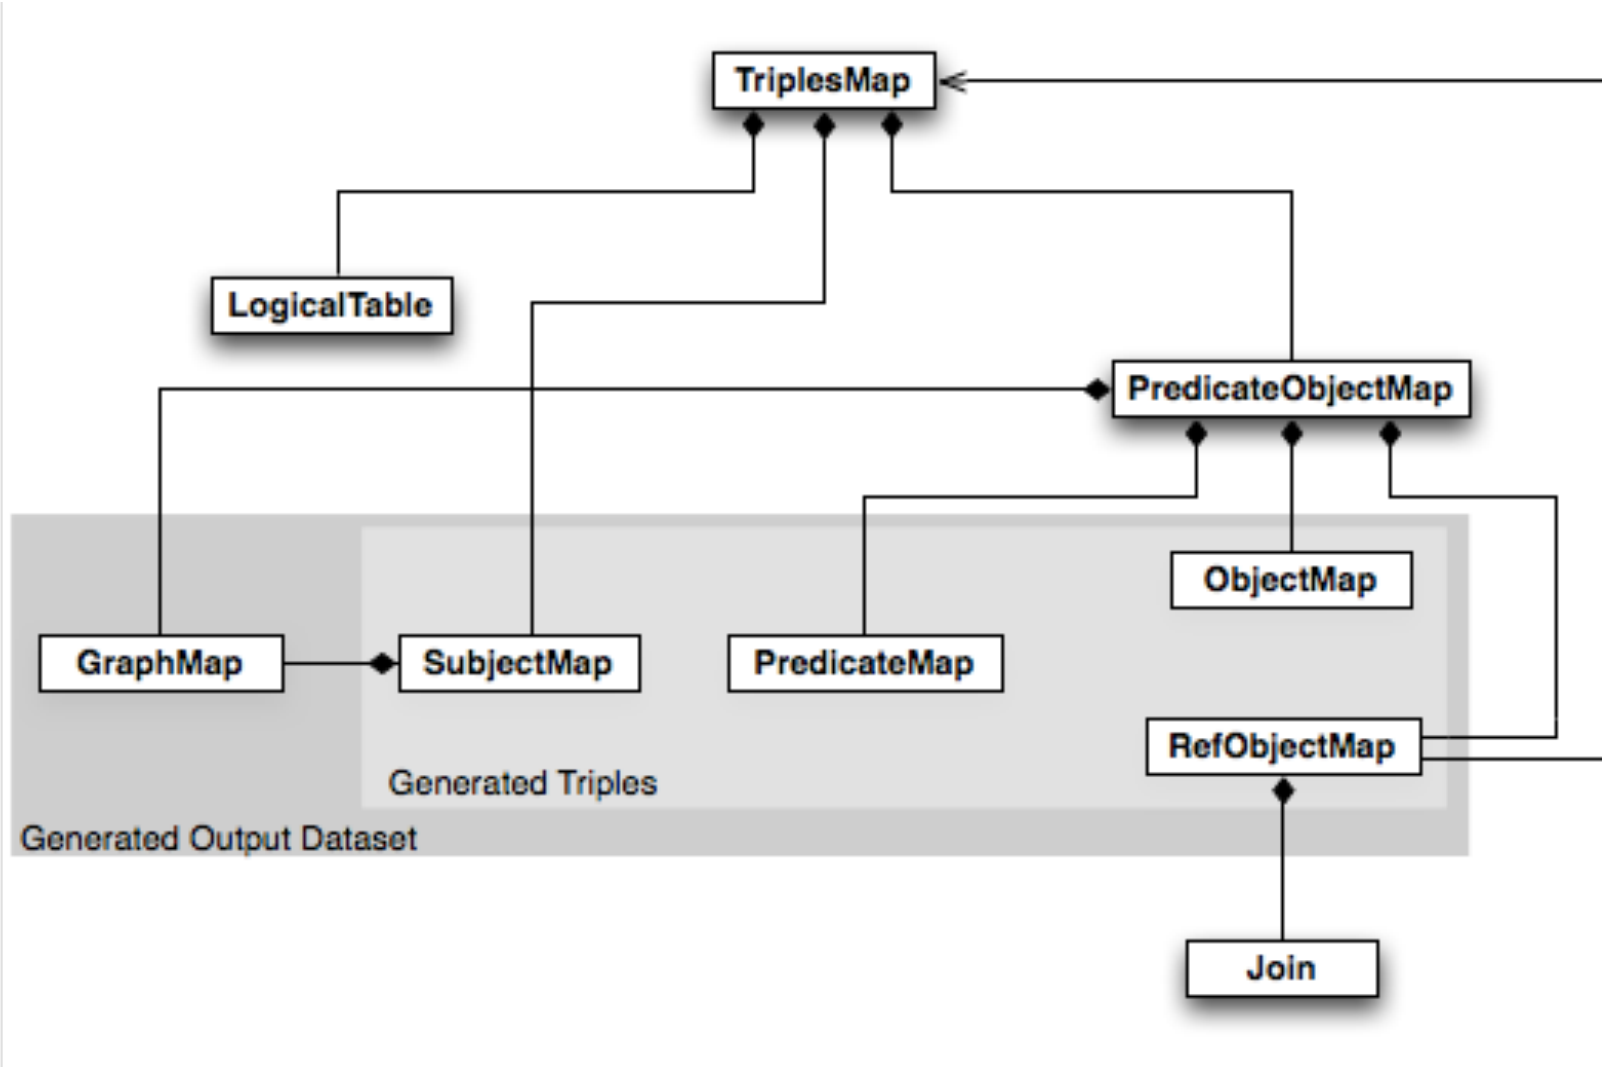
R2RML: Running Example



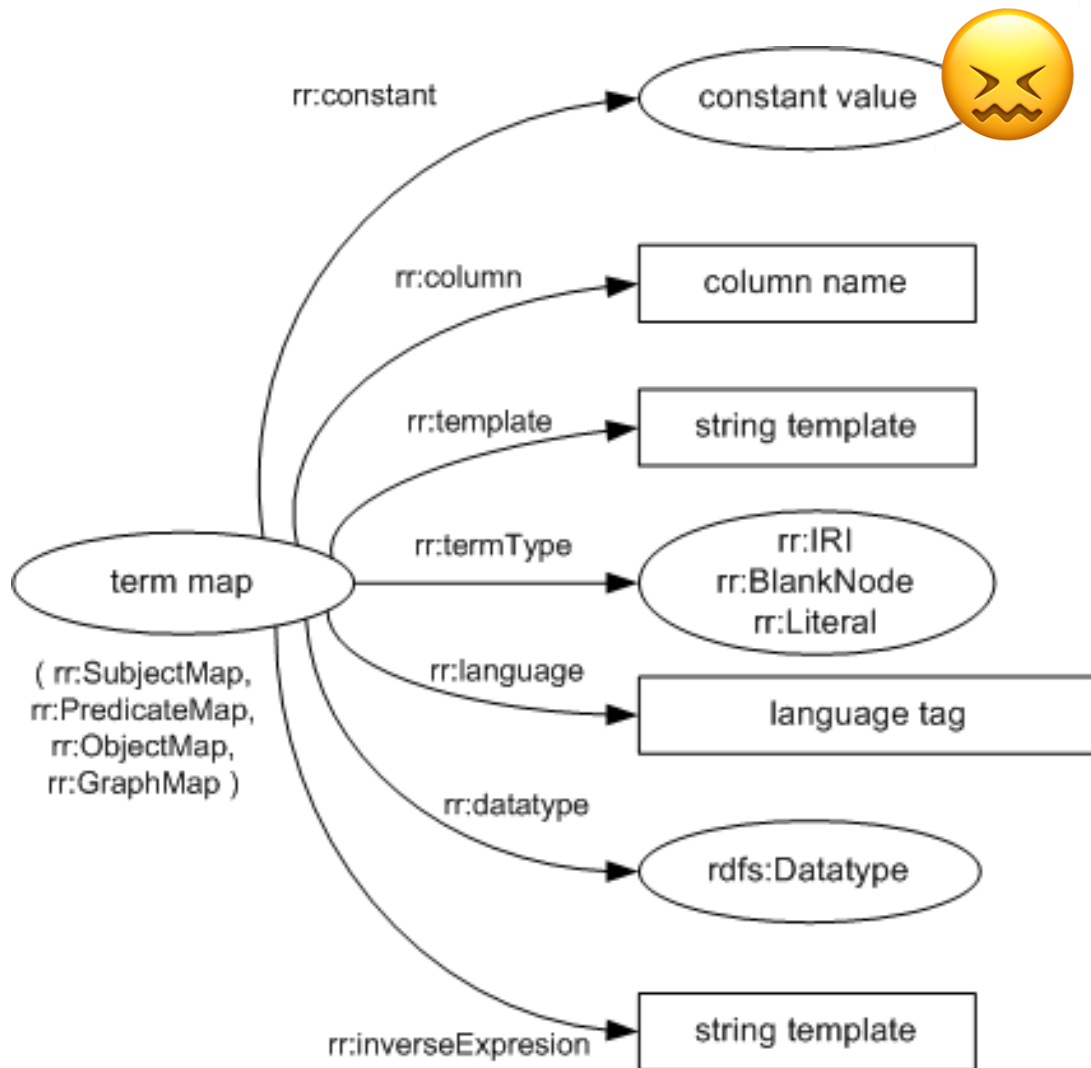
The screenshot shows a text editor window titled "output.ttl — r2rml-ois" with a standard macOS window header (red, yellow, green buttons) and an "Add License" button. The editor contains 17 lines of R2RML Turtle code. A vertical line number column on the left shows lines 1 through 17. The code defines three R2RML patterns for Person and Address resources. The first pattern (lines 1-5) is for a person named Kevin. The second pattern (lines 6-12) is for a person named Christophe. The third pattern (lines 13-16) is for an address named Brussels. The code uses URIs for the classes and properties, and literals for the values. The editor has a status bar at the bottom showing "Line: 1", "Turtle" mode, and "Tab Size: 4".

```
1 <http://foo.example/Person/2>
2   a      <http://xmlns.com/foaf/0.1/Person> ;
3   <http://xmlns.com/foaf/0.1/name>
4     "Kevin" .
5
6 <http://foo.example/Person/1>
7   a      <http://xmlns.com/foaf/0.1/Person> ;
8   <http://xmlns.com/foaf/0.1/based_near>
9     <http://foo.example/Addresses/1> ;
10  <http://xmlns.com/foaf/0.1/name>
11    "Christophe" .
12
13 <http://foo.example/Addresses/1>
14   a      <http://dbpedia.org/ontology/Place> ;
15   <http://xmlns.com/foaf/0.1/name>
16     "Brussels" .
17
```

R2RML



R2RML



- If `rr:constant` is used for a Subject Map, Graph Map, or Predicate Map, then the value must be an IRI.
- If `rr:constant` is used for an Object Map, then its value must be either a literal or an IRI.
- So either a rectangle or an ellipse, depending on which...
- `rr:inverseExpressions` are useful only when accessing RDBs as virtual graphs (they are an "optimization hint").

rr:class in rr:SubjectMaps

```
rr:subjectMap [  
  rr:template "http://foo.example/Addresses/{ID}" ;  
  rr:class dbpedia:Place  
] ;
```

This subject map generates for each subject a triple with `rdf:type` as predicate and `dbpedia:Place` as object. In other words, `rdf:type` and `dbpedia:Place` are constants. Can we use an `rr:PredicateObjectMap` with two constants?

Yes:

```
rr:predicateObjectMap [  
  rr:predicate rdf:type ;  
  rr:object dbpedia:Place  
] ;
```

When the FK is in your table

People		
PK		→Addresses(ID)
ID	fname	addr

Addresses	
PK	
ID	city

When you join tables based on a FK, and the PK of the parent is used for the Subject Map, there can be a possibility to avoid join. How?

```
<#PersonTriplesMap>
a rr:TriplesMap;
rr:logicalTable [ rr:tableName "People" ] ;
rr:subjectMap [
  rr:template "http://foo.example/Person/{ID}" ;
  rr:class foaf:Person
] ;
rr:predicateObjectMap [
  rr:predicate foaf:name ;
  rr:objectMap [ rr:column "fname" ]
] ;
rr:predicateObjectMap [
  rr:predicate foaf:based_near ;
  rr:objectMap [
    rr:parentTriplesMap <#AddressTripleMap> ;
    rr:joinCondition [
      rr:child "addr" ;
      rr:parent "ID"
    ]
  ]
] ;
.
```

When the FK is in your table

People		
PK		→Addresses(ID)
ID	fname	addr

Addresses	
PK	
ID	city

Create a Term Map that generates the same resource!

Prone to errors if, for instance, templates change, though.

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ]
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:template "http://foo.example/Addresses/{addr}" ;
    ]
  ] ;
  .
```

Graph Maps

- Any subject map or predicate-object map may have one or more associated graph maps. They are specified in one of two ways. Either by associating a constant with the property `rr:graph`, or by providing a Graph Map (which returns IRIs) with the property `rr:graphMap`. `rr:defaultGraph` is reserved for the default (nameless) graph.
- If a Subject Map has no Graph Maps then the set of Graph Maps is `{rr:defaultGraph}`.
- If both the Subject Map and a Predicate Object Map have no Graph Maps, then the set of Graph Maps is `{rr:defaultGraph}`. Otherwise it is, for each Predicate Object Map, the union of both graph sets.
- It is actually more simple than it sounds, let's exemplify!

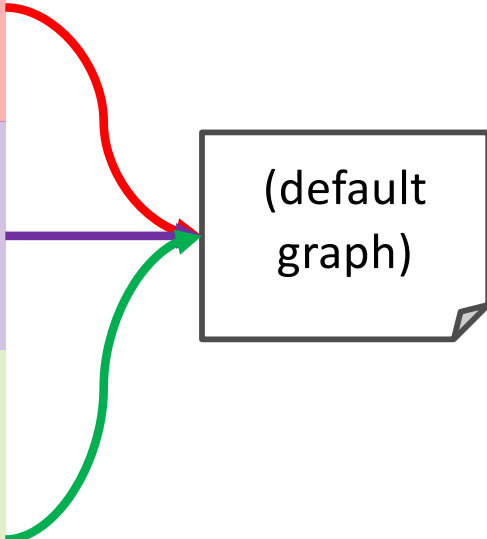
Graph Maps

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person ;

  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ] ;

  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [ rr:child "addr" ; rr:parent "ID" ]
    ] ;

  ] ;
.
```



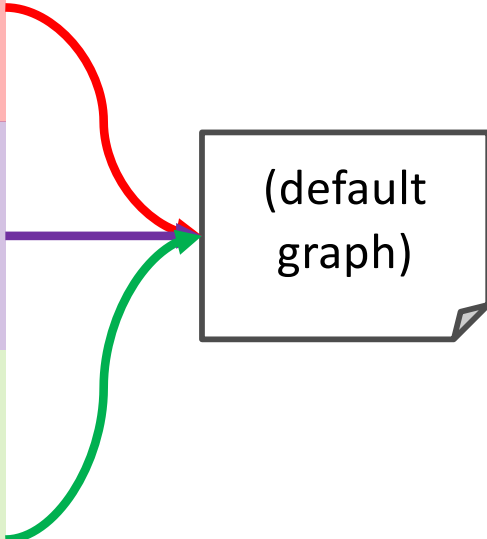
(default graph)

Graph Maps

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person ;
    rr:graph rr:defaultGraph
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ] ;

  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [ rr:child "addr" ; rr:parent "ID" ]
    ] ;

  ] ;
.
```



(default graph)

Graph Maps

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person ;
    rr:graph <#one>
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ] ;

  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [ rr:child "addr" ; rr:parent "ID" ]
    ] ;

  ] ;
```



<#one>

Graph Maps

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person ;
    rr:graph <#one>
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ] ;
    rr:graph <#two>
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [ rr:child "addr" ; rr:parent "ID" ]
    ] ;
    rr:graph <#three>
  ] ;
.
```

<#one>

<#two>

<#three>

Graph Maps

```
<#PersonTriplesMap>
  a rr:TriplesMap;
  rr:logicalTable [ rr:tableName "People" ] ;
  rr:subjectMap [
    rr:template "http://foo.example/Person/{ID}" ;
    rr:class foaf:Person ;
    rr:graph <#one>
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:name ;
    rr:objectMap [ rr:column "fname" ] ;
    rr:graph <#two>
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:based_near ;
    rr:objectMap [
      rr:parentTriplesMap <#AddressTripleMap> ;
      rr:joinCondition [ rr:child "addr" ; rr:parent "ID" ]
    ] ;
    rr:graph <#three>, rr:defaultGraph
  ] ;
```

<#one>

<#two>

(default
graph)

<#three>

R2RML

- R2RML supports mapping values with constants, column values or column values applied to a template.
- If an objectMap does not refer to a column or has no language tag, then term types default to IRIs unless you explicitly specify it to be a Literal.

...

```
rr:predicateObjectMap [  
  rr:predicateMap [ rr:constant foaf:name ] ;  
  rr:objectMap [  
    rr:template "{first} {last}" ;  
    rr:termType rr:Literal ;  
  ]  
]
```

...

R2RML

- Term maps with a TermType of rr:Literal can have a language tag.

...

```
rr:predicateObjectMap [  
  rr:predicateMap [ rr:constant rdfs:label ] ;  
  rr:objectMap [  
    rr:column "{title}" ;  
    rr:language "en" ;  
  ]  
]
```

...

- Why is the above TermType implied to be a rr:Literal? It uses a column.

R2RML: Multiple Languages

What if you have a table with multiple languages?

Assuming you have a discriminator, you can

- Create a logical table with an SQL query for each language.
 - How?
- Create one logical table and use a language column to create a single mapping for all languages.
 - Unfortunately, this is not part of the recommendation and support depends on the implementation...

```
rr:objectMap [  
  rr:column "TITLE" ;  
  rrx:languageColumn "TITLE_LANG" ;  
].
```

R2RML: Datatypes

Datatypes can only be declared for TermMaps that are of type `rr:Literal` and without a language tag.

```
rr:objectMap [  
  rr:column "EMPNO" ;  
  rr:datatype xsd:positiveInteger  
]
```

R2RML

- R2RML provides a highly customizable language for mapping relational databases to triples.
- Unlike direct mapping that reflects the database's structure, the author of the mapping decides on the structure and ontology.
 - Mapping *projections*, e.g., a tripleMap for $\pi_{\{city\}}(Addresses)$ to create instances of the concept `ex:City` (cfr. “degroupping” from lecture 1)
 - Mapping *selections*, e.g., a tripleMap for $\sigma_{\{type='Cat'\}}(Pet)$ to create instances of `ex:Cat`
 - ...

rr:inverseExpression

- Interesting for **virtual graphs** – retrieving RDF by translating SPARQL queries into SQL queries via a mediator.
- R2RML transforms “database terms” into RDF terms. Inverse expressions help one to transform RDF terms back into “database terms” which may increase performance.
- I am not aware of any implementation that supports or implements this, as it is not necessary for R2RML views to work.

rr:inverseExpression

R2RML: we assume surnames are stored in capital letters

```
<TriplesMap1> a rr:TriplesMap ;
  rr:logicalTable [
    rr:sqlQuery "SELECT email, LOWER(surname) AS lsur FROM emp" ; ] ;
  rr:subjectMap [
    rr:template "http://example.org/person/{lsur}" ;

  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:mbox ;
    rr:objectMap [ rr:column "email" ]
  ] .
```

The SQL query for retrieving the value of an email address should correspond with:
SELECT email FROM emp WHERE
LOWER(surname) = 'debruyne'
This may be slow

SPARQL:

```
SELECT ?email {
  <http://example.org/person/debruyne> foaf:mbox ?email .
}
```


rr:inverseExpression

R2RML: we assume surnames are stored in capital letters

```
<TriplesMap1> a rr:TriplesMap ;
  rr:logicalTable [
    rr:sqlQuery "SELECT email, LOWER(surname) AS lsur FROM emp" ; ] ;
  rr:subjectMap [
    rr:template "http://example.org/person/{lsur}" ;
    rr:inverseExpression "{surname} = UPPER({lsur})" ;
  ] ;
  rr:predicateObjectMap [
    rr:predicate foaf:mbox ;
    rr:objectMap [ rr:column "email" ]
  ] .
```

We can now optimize the query to
SELECT email FROM emp
WHERE **surname = UPPER('debruyne')**
If surname were to be indexed, this query
would be much faster!

SPARQL:

```
SELECT ?email {
  <http://example.org/person/debruyne> foaf:mbox ?email .
}
```

RDB2RDF

- Translate data to RDF: generates an RDF dump for immediate consumption, but will be harder to maintain (A)
 - Loaded into triplestores such as Jena TDB (with Fuseki for a SPARQL endpoint), or Virtuoso.
- A mapping from the RDB to RDF, generating SPARQL queries into (intermediate) SQL queries, but will have longer query times (B)
- The TA will demonstrate how you can use OnTop to mediate between ontologies and a database via mappings.



R2RML-F

- Introduction of a *Function Valued* Term Map that allow for user defined functions at the cost of tractability.
- C. Debruyne and D. O'Sullivan. R2RML-F: Towards Sharing and Executing Domain Logic in R2RML Mappings. In Proceedings of the Workshop on Linked Data on the Web, LDOW 2016, co-located with the 25th International World Wide Web Conference (WWW 2016), Montreal, Canada, April 12th, 2016, 2016.
- <https://github.com/chrdebru/r2rml>

Extending R2RML

- Namespace **rdf**: <http://kdeg.scss.tcd.ie/ns/rdf#>
- *Functions* have a function *name* and *body*.

```
<#Multiply>
  rdf:functionName "multiply" ;
  rdf:functionBody """
    function multiply(var1, var2) {
      return var1 * var2 ;
    }
  """ ;
.
```

- Functions are written in **ECMAScript**.

Extending R2RML

A “function valued” term map *calls a function* and the *parameters* are themselves term maps.

```
<#TriplesMap1>
rr:logicalTable [ rr:tableName "Employee"; ];
rr:subjectMap [ rr:template "http://org.com/employee/{ID}"; ] ;
rr:predicateObjectMap [
  rr:predicate ex:salary ;
  rr:objectMap [
    rr:datatype xsd:double ;
    rrf:functionCall [
      rrf:function <#Multiply> ;
      rrf:parameterBindings (
        [ rr:constant "12"^^xsd:integer ]
        [ rr:column "monthly_salary" ]
      ) ;
    ] ;
  ] ;
]
```

Parameter bindings as
an RDF Collection.

Parameter bindings can
be empty.

Term Maps as parameters.

RML

- Developed by imec at UGent
 - <http://semweb.mmlab.be/rml/spec.html>
 - <https://github.com/mmlab/RMLProcessor>
- An extension of R2RML to support XML, HTML, CSV, JSON...
 - “Basically” a superset of R2RML.
- See also: A. Dimou, M. Vander Sande, P. Colpaert, R. Verborgh, E. Mannens, R. Van De Walle: “RML: A Generic Language for Integrated RDF Mappings of Heterogeneous Data”, in Proceedings of the 7th Workshop on Linked Data on the Web, WWW14, 2014 Seoul, Korea.

Other tools and mapping languages

- Stardog -- <http://stardog.com/>
- D2RQ -- <http://d2rq.org/>
- Virtuoso -- <http://virtuoso.openlinksw.com/>

An interactive demonstration

Weatherstations.csv – Weather Readings omitted in this slide.

Name	Weather_Reading	Agency	LAT	LONG
M50 Blanchardstown	...	National Roads Authority	53.37046603	-6.380851447
M50 Dublin Airport	...	National Roads Authority	53.40964111	-6.227597428
Dublin Airport	...	Met ...ireann	53.42150608	-6.29784754

Using the GeoSPARQL, we will create instances of geo:Feature (the weather stations) and their instances of geo:Geometry (a point using the latitude and longitude).

References

- Satya S. Sahoo et al. A Survey of Current Approaches for Mapping of Relational Databases to RDF. W3C RDB2RDF XG Report, W3C, 2009. http://www.w3.org/2005/Incubator/rdb2rdf/RDB2RDF_SurveyReport.pdf
- Berners-Lee, T. Relational Databases on the Semantic Web, 1998, via <http://www.w3.org/DesignIssues/RDB-RDF.html>
- R. Cyganiak, C. Bizer, J. Garbers, O. Maresch, and C. Becker. The D2RQ mapping language. <http://d2rq.org/d2rq-language>, March 2012.
- E. Prud'hommeaux and A. Seaborne. SPARQL Query Language for RDF. W3C Recommendation, W3C, January 2008.